Dataproc

Dataproc

- Dataproc is a managed Apache Spark and Apache Hadoop service that lets you take advantage of open-source data tools for batch processing, querying, streaming, and machine learning.
- Dataproc automation helps you create clusters quickly, manage them easily, and save money by turning clusters off when you don't need them.
- With less time and money spent on administration, you can focus on your jobs and your data.

Advantages of Dataproc

Low Cost

- Dataproc is priced at only 1 cent per virtual CPU in your cluster per hour, on top of the other Cloud Platform resources you use.
- Dataproc clusters can include preemptible instances that have lower compute prices, reducing your costs even further.
- Instead of rounding your usage up to the nearest hour, Dataproc charges you only for what you really use with second-by-second billing.

Advantages of Dataproc

Super Fast

- Dataproc clusters are quick to start, scale, and shutdown, with each of these operations taking 90 seconds or less, on average.
- Spend less time waiting for clusters and more hands-on time working with your data.

Advantages of Dataproc

Integrated

- Dataproc has built-in integration with other Google Cloud Platform services, such as BigQuery, Cloud Storage, Cloud Bigtable, Cloud Logging, and Cloud Monitoring
- It is more than just a Spark or Hadoop cluster - a complete data platform.
- Ex. Dataproc can be used to effortlessly ETL terabytes of raw log data directly into BigQuery for business reporting.

Advantages of Dataproc

Managed

- Use Spark and Hadoop clusters without needing an administrator or special software
- Easily interact with clusters and jobs through Google Cloud console, Cloud SDK, or Dataproc REST API
- Turn off clusters when not in use to avoid costs on idle resources
- Data remains safe as Dataproc integrates with Cloud Storage, BigQuery, and Cloud Bigtable

Advantages of Dataproc

Simple

- No need to learn new tools or APIs; existing projects move to Dataproc without redevelopment
- Spark, Hadoop, Pig, and Hive are frequently updated for faster productivity

Dataproc Cluster Types

- Standard (1 master, N workers)
- Single Node (1 master, 0 workers)
- High Availability (3 masters, N workers)

Dataproc Cluster Types

- Jobs Supported
  - Hadoop
  - Spark
  - SparkR
  - PySpark
  - SparkSQL
  - Hive

Cluster Activities

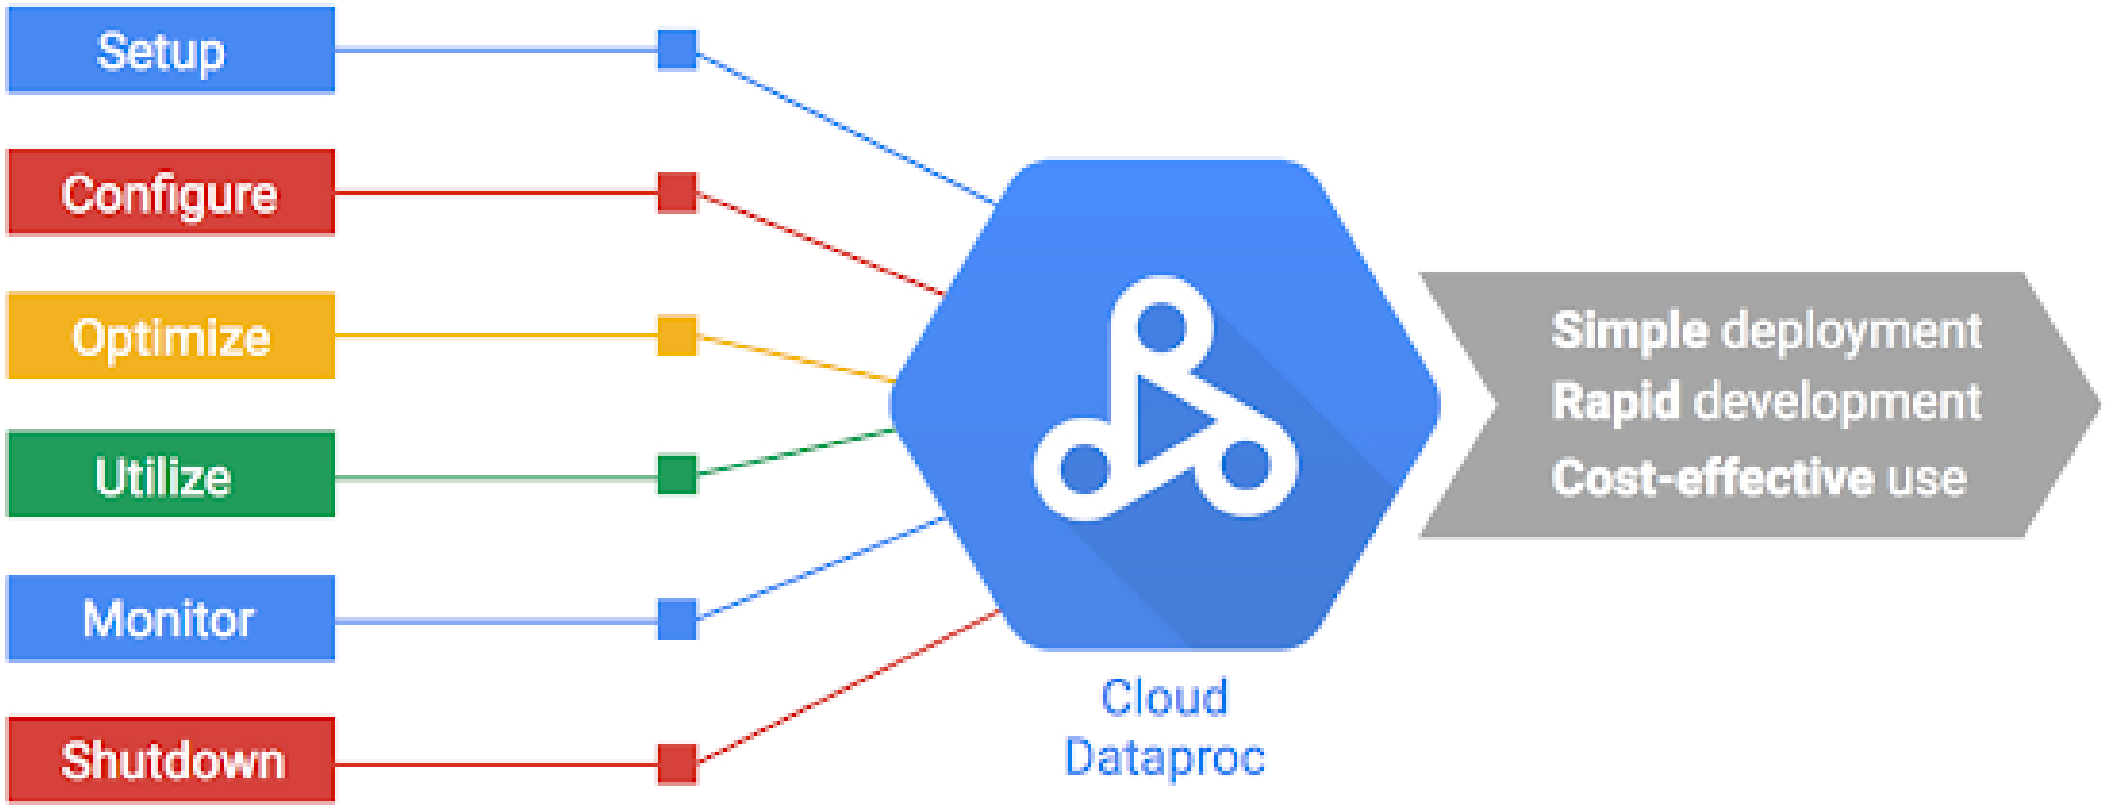

Cluster Naming Rules

- A cluster name must start with a lowercase letter.
- Up to 51 characters allowed: lowercase letters, numbers, and hyphens.
- A name cannot end with a hyphen.

Cluster Naming Rules

- A cluster name must start with a lowercase letter.
- Up to 51 characters allowed: lowercase letters, numbers, and hyphens.
- A name cannot end with a hyphen.

Cluster Region

- A Compute Engine region must be specified for the cluster (e.g., us-east1, europe-west1).
- Regional isolation ensures that VM instances, metadata, and Cloud Storage remain within the selected region.
- Region details can be viewed with the gcloud compute regions list command.
- Within a region, resources are further divided into zones (e.g., us-east1-b).
- Zones provide redundancy and fault isolation for cluster components.
- Cluster nodes are typically distributed across zones for high availability.

Cluster Connectivity

- A Dataproc cluster requires full internal IP networking cross-connectivity between master and worker VMs.
- The default VPC network provides this connectivity automatically.
- Networking can be customized for advanced configurations such as private clusters.

Master and Worker Nodes

- Master node manages cluster metadata, scheduling, and monitoring.
- Worker nodes run the actual distributed data processing tasks.
- Additional secondary workers can be added for preemptible, cost-saving compute capacity.

Machine Types and Scaling

- Each node in a cluster is assigned a Compute Engine machine type (e.g., c4-standard-4).
- Machine types define CPU, memory, and storage capacity.
- Clusters can scale by adding or removing workers dynamically.

Storage Integration

- Dataproc integrates with Cloud Storage for storing input, output, and temporary data.
- Cluster data persists in Cloud Storage, even if the cluster is deleted.
- HDFS is supported but typically used only for temporary storage.

Security and IAM

- Cluster access and management are secured with Identity and Access Management (IAM).
- Encryption of data at rest and in transit is enabled by default.
- Service accounts control permissions for clusters accessing other Google Cloud services.

Cost and Lifecycle Management

- Clusters can be turned off when not needed, saving costs on idle resources.
- Preemptible VMs lower compute costs.
- Auto-deletion policies can be configured to shut down clusters after jobs are completed.

Preemptible VMs

- **Short-Lived Instances:** Virtual machines that can run for a maximum of 24 hours.
- **Cost Savings:** Cheaper than standard VMs.
- **Cluster Use:** Often added as secondary worker nodes in Dataproc clusters to handle extra processing.
- **Preemption:** Can be stopped by Google Cloud at any time if resources are needed elsewhere.
- Best For Batch jobs, big data analytics, and machine learning workloads that can tolerate interruptions.

HDFS Integration in Dataproc

- Dataproc integrates with Apache Hadoop and uses the Hadoop Distributed File System (HDFS) for on-cluster storage.
- Additionally, Dataproc installs an HDFS-compatible Cloud Storage connector, enabling Blob storage (GCS) to be used in parallel with HDFS.
- Data can be uploaded to or downloaded from clusters using either HDFS or Cloud Storage.
- VM boot disks (hosting HDFS data) are tied to cluster lifetime: they are deleted when the Dataproc cluster is deleted.
- To preserve data beyond the cluster lifespan, it is recommended to store files in Cloud Storage.

Scaling a Dataproc Cluster

- Scaling a Dataproc cluster involves increasing or decreasing the number of worker nodes—both primary and secondary.
- This change can be made even while jobs are running.
- Horizontal scaling adjusts the number of worker nodes.
- Vertical scaling (changing machine types, CPUs, memory) is not supported once the cluster is created.
- Scaling applies to primary worker nodes or secondary (e.g., preemptible) workers, or both.
- New nodes must match the machine type of existing workers.

Autoscaling for Smart Scaling

- Rather than manual updates, Autoscaling can automatically adjust worker counts.
- Autoscaling uses preconfigured policies to adjust resource levels based on workload.

Create a Dataproc Cluster

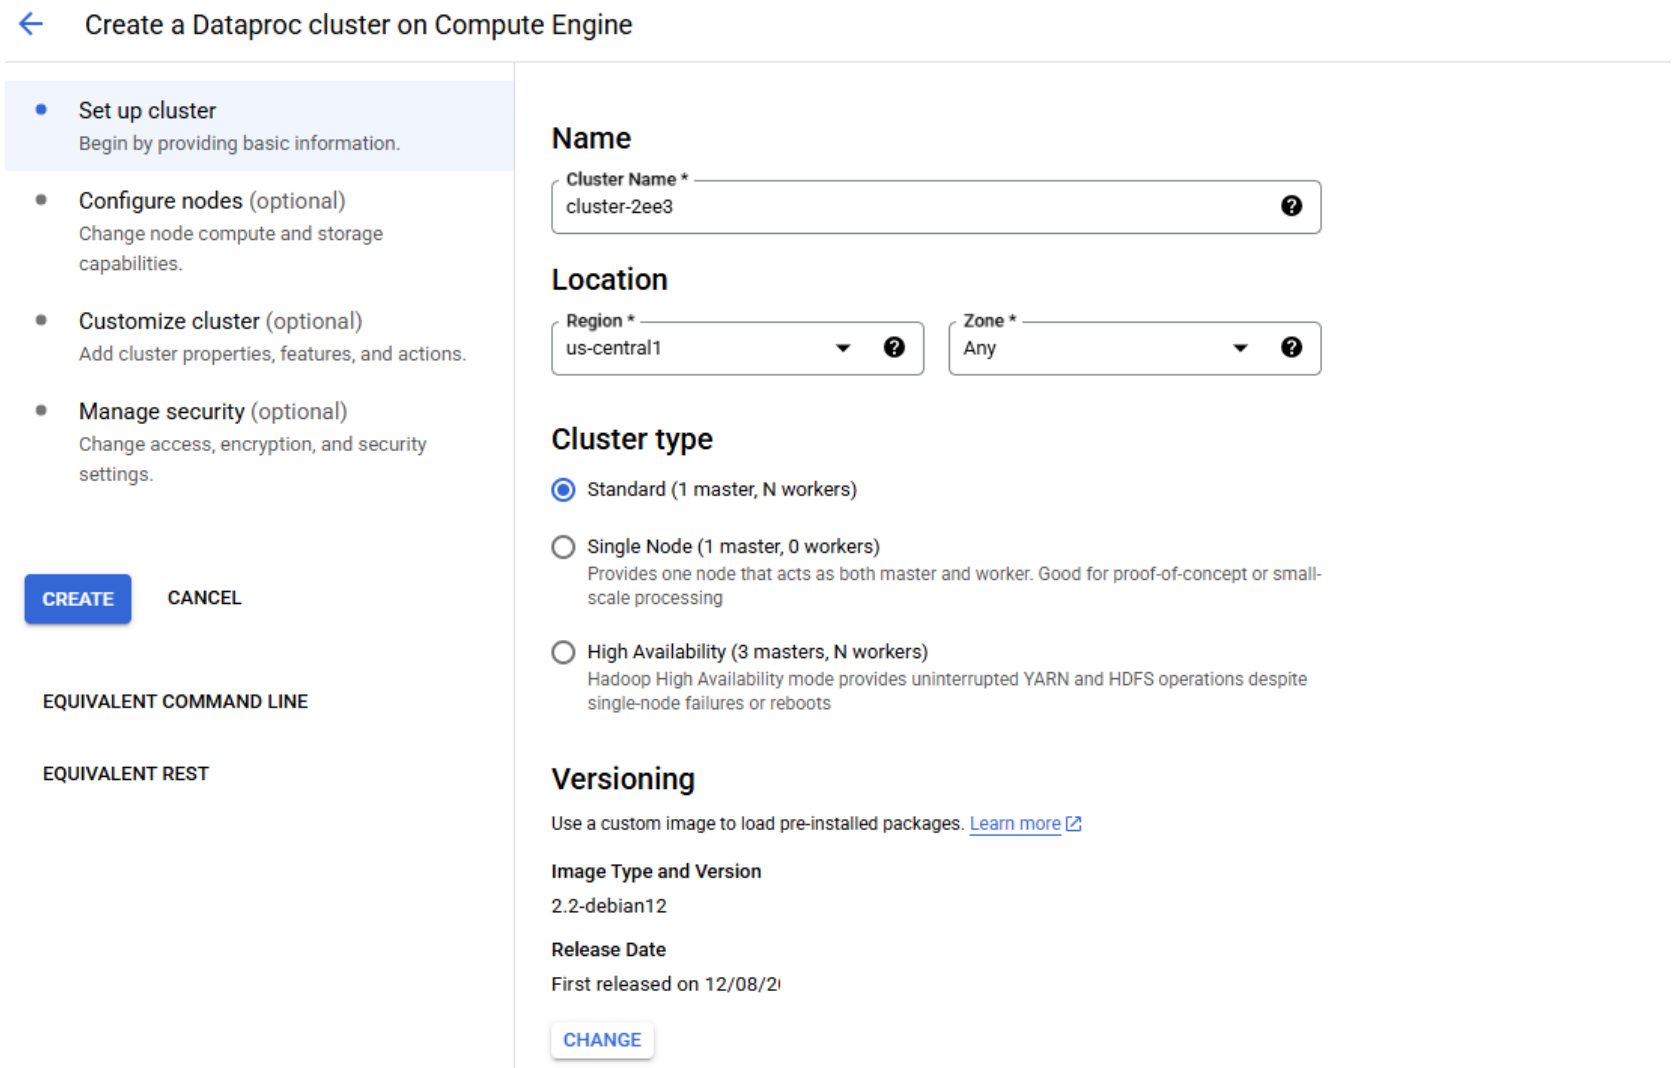

Create a Dataproc Cluster

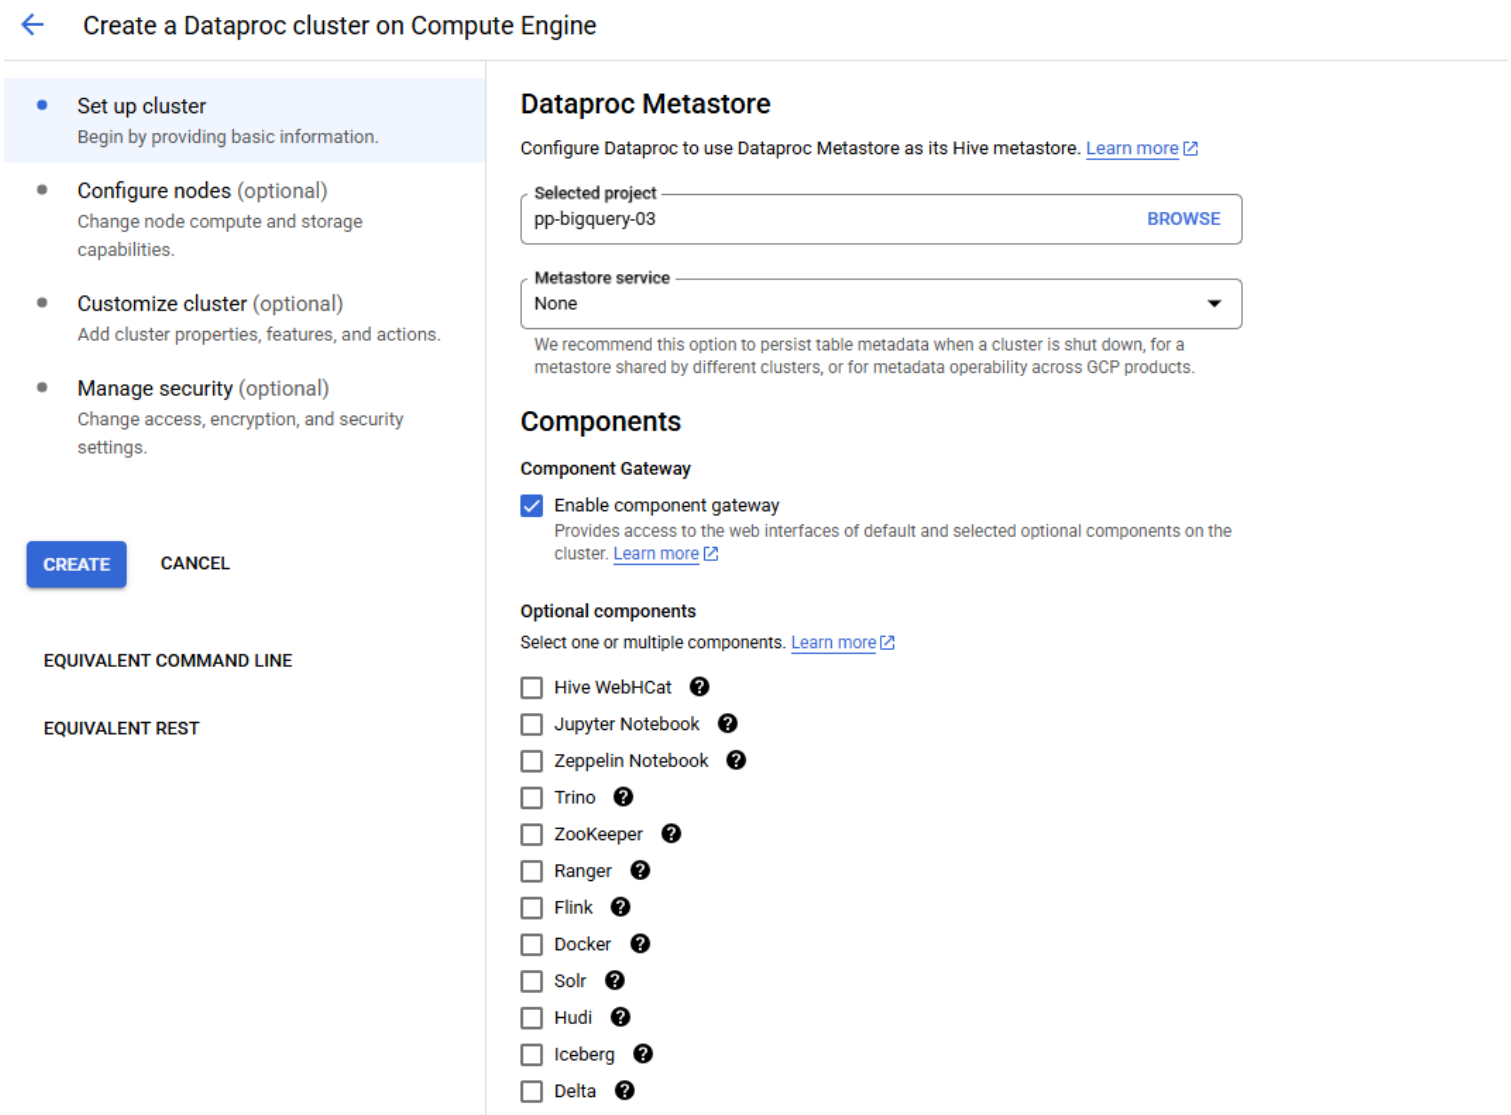

Create a Dataproc Cluster

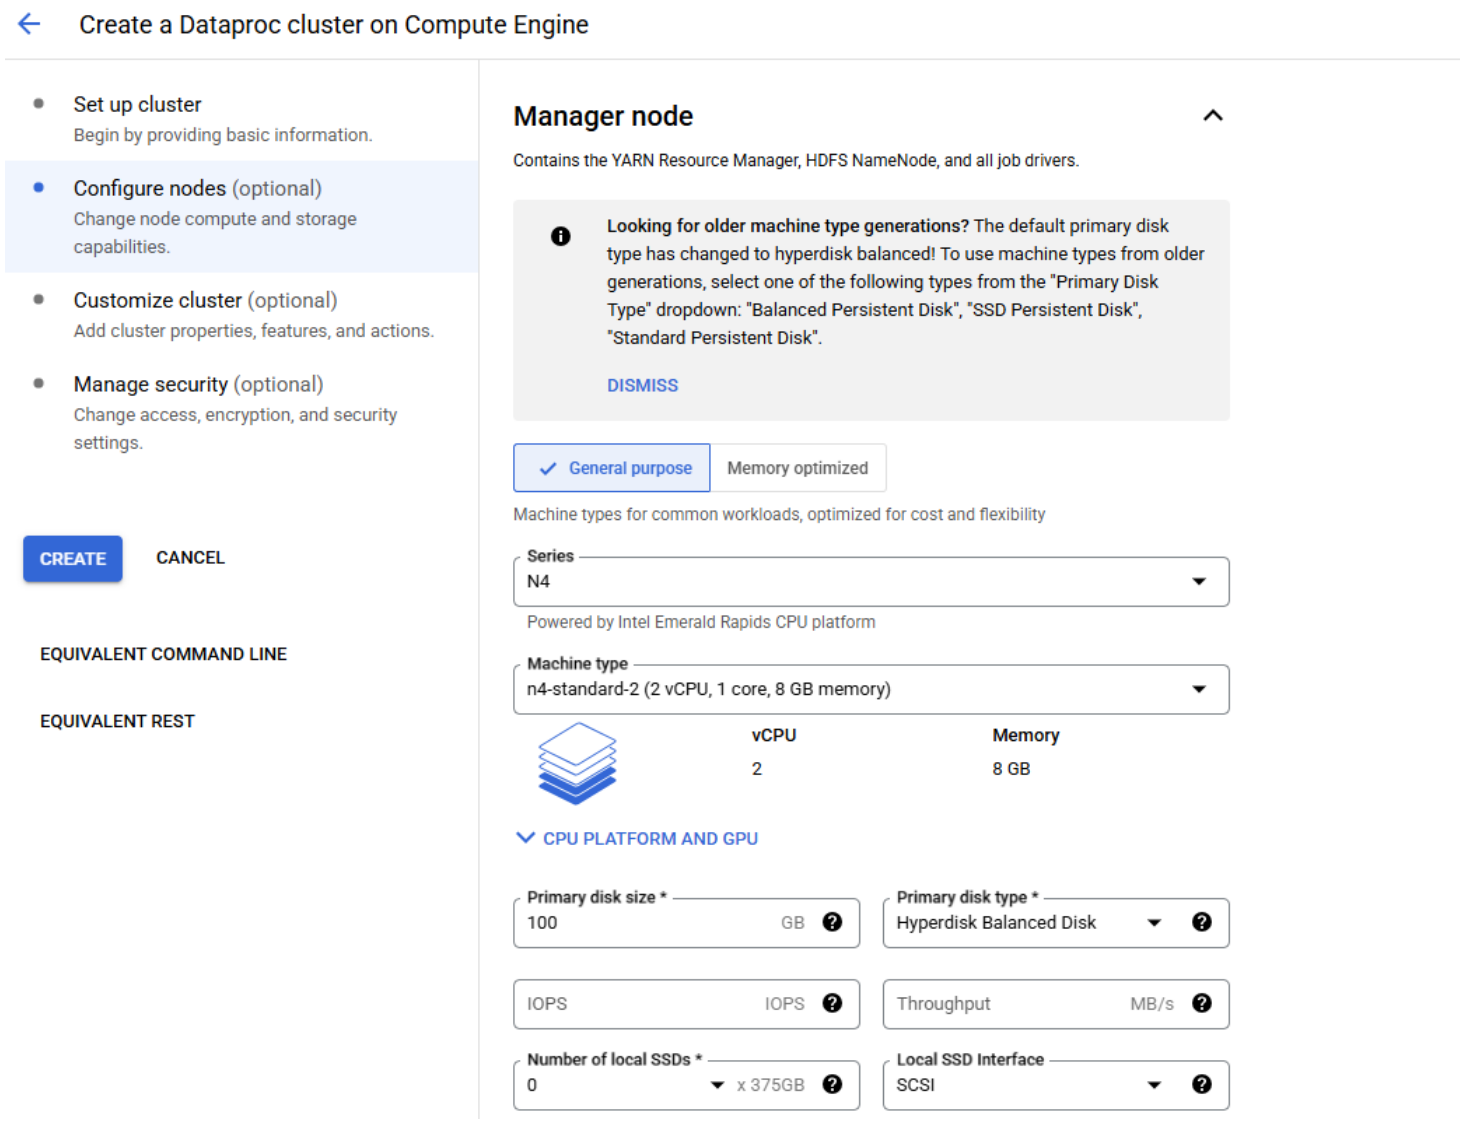

Create a Dataproc Cluster

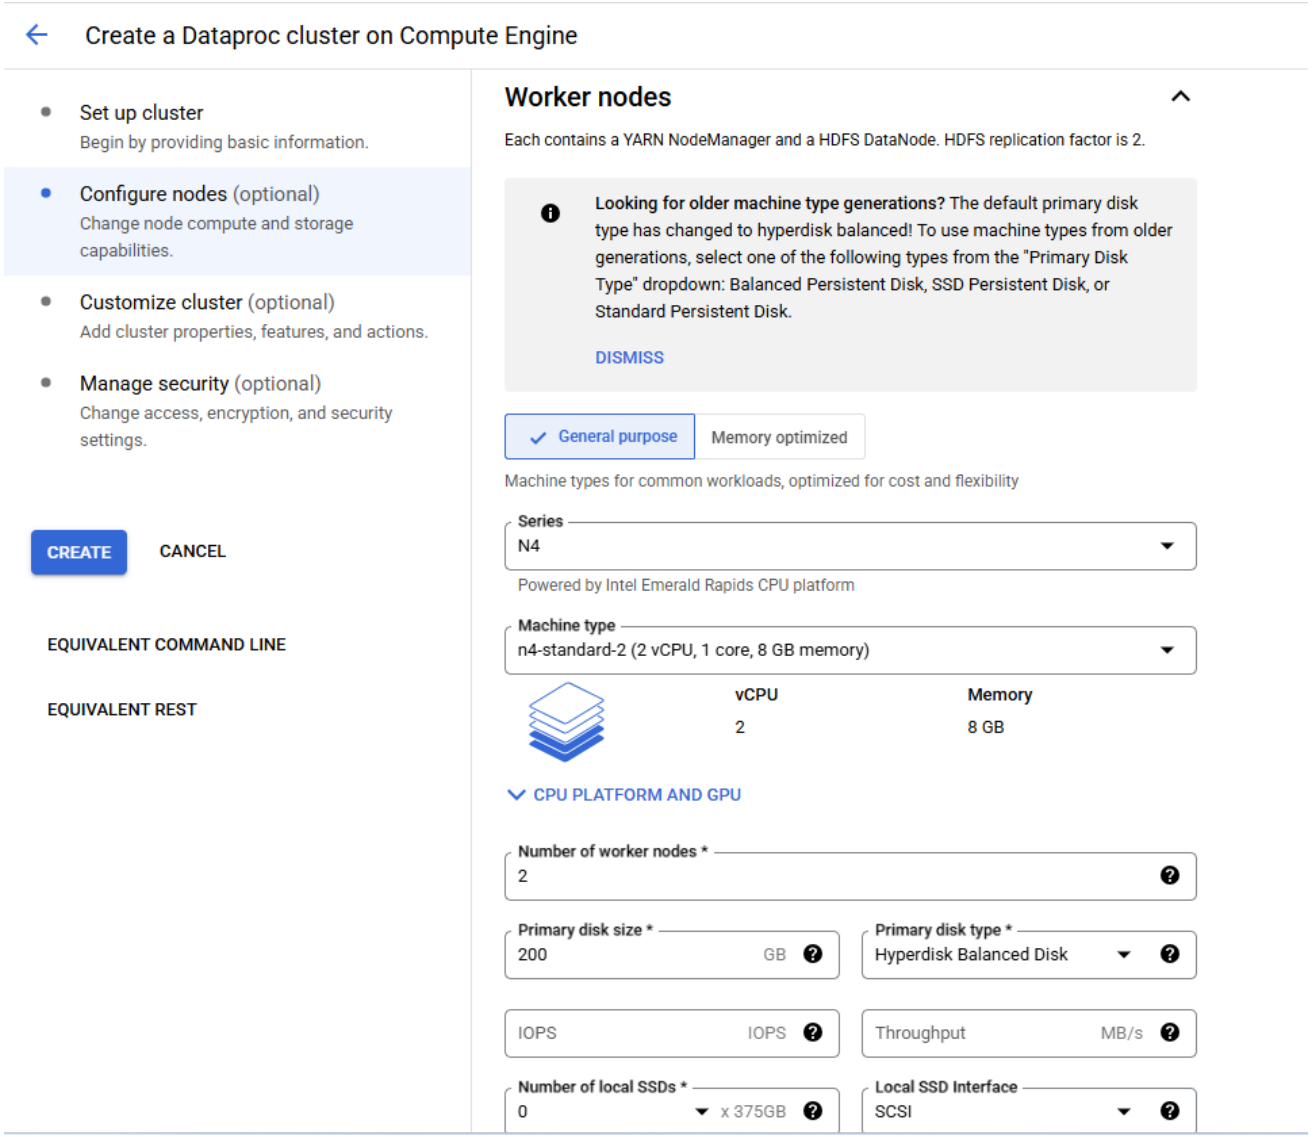

Create a Dataproc Cluster

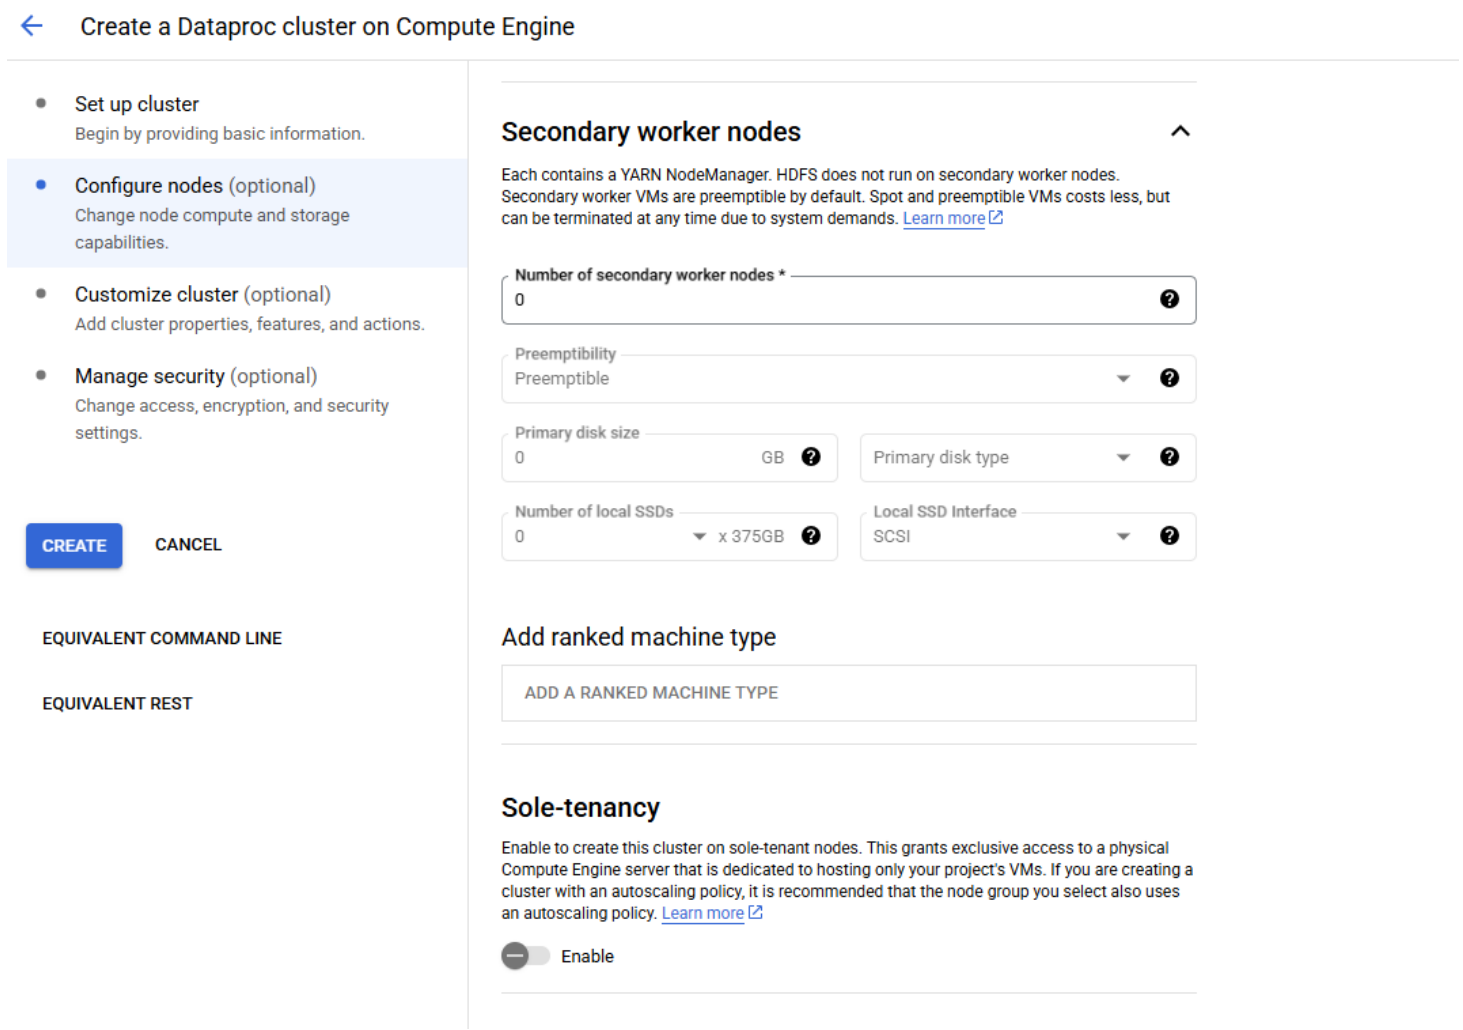

Create a Dataproc Cluster

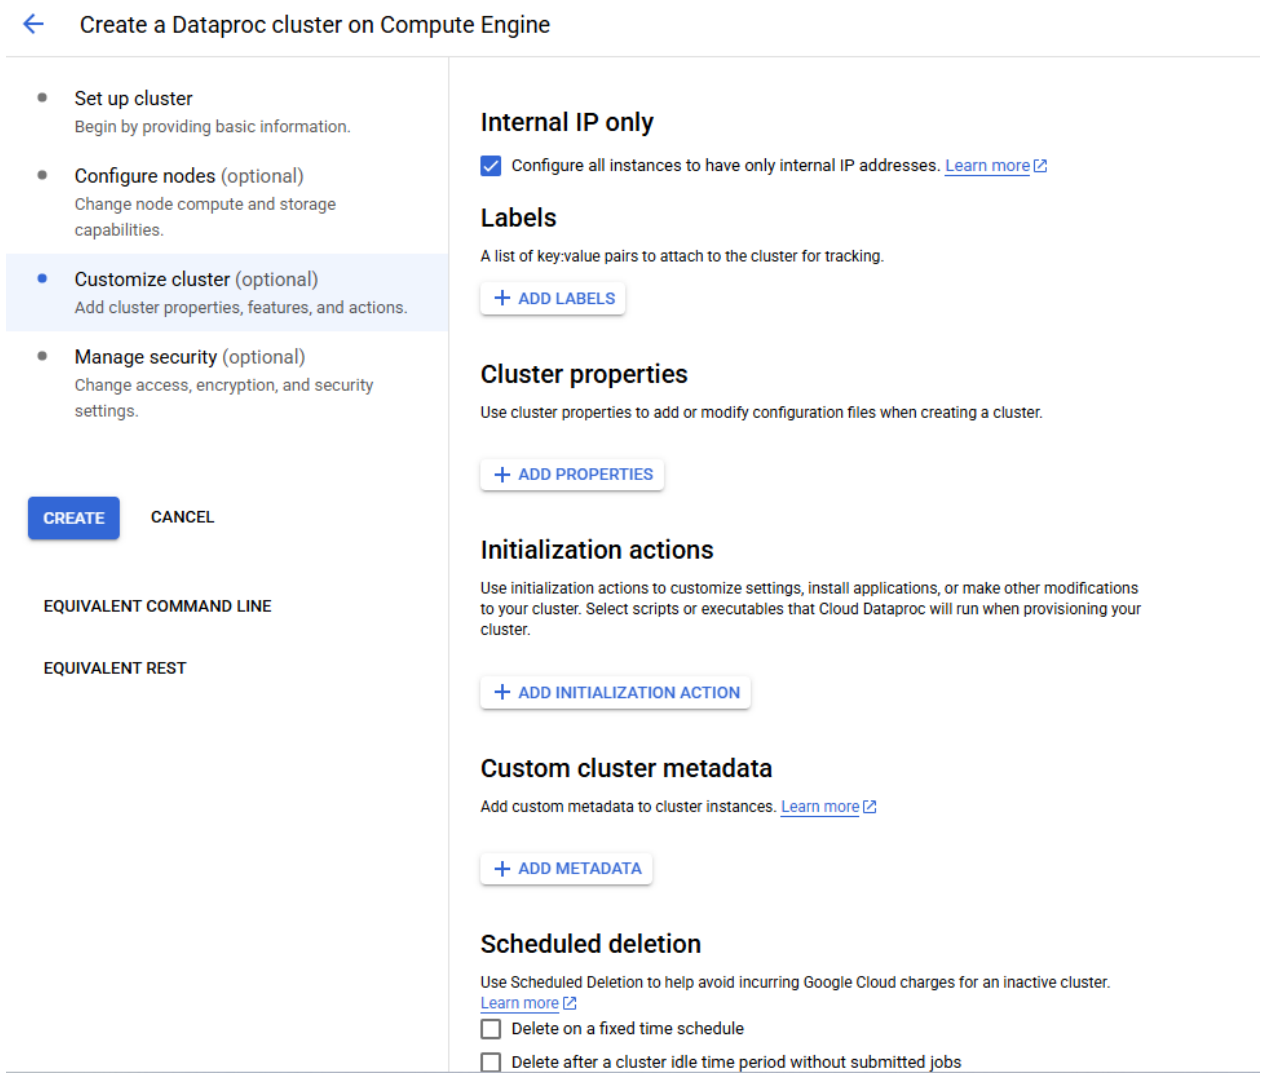

Create a Dataproc Cluster

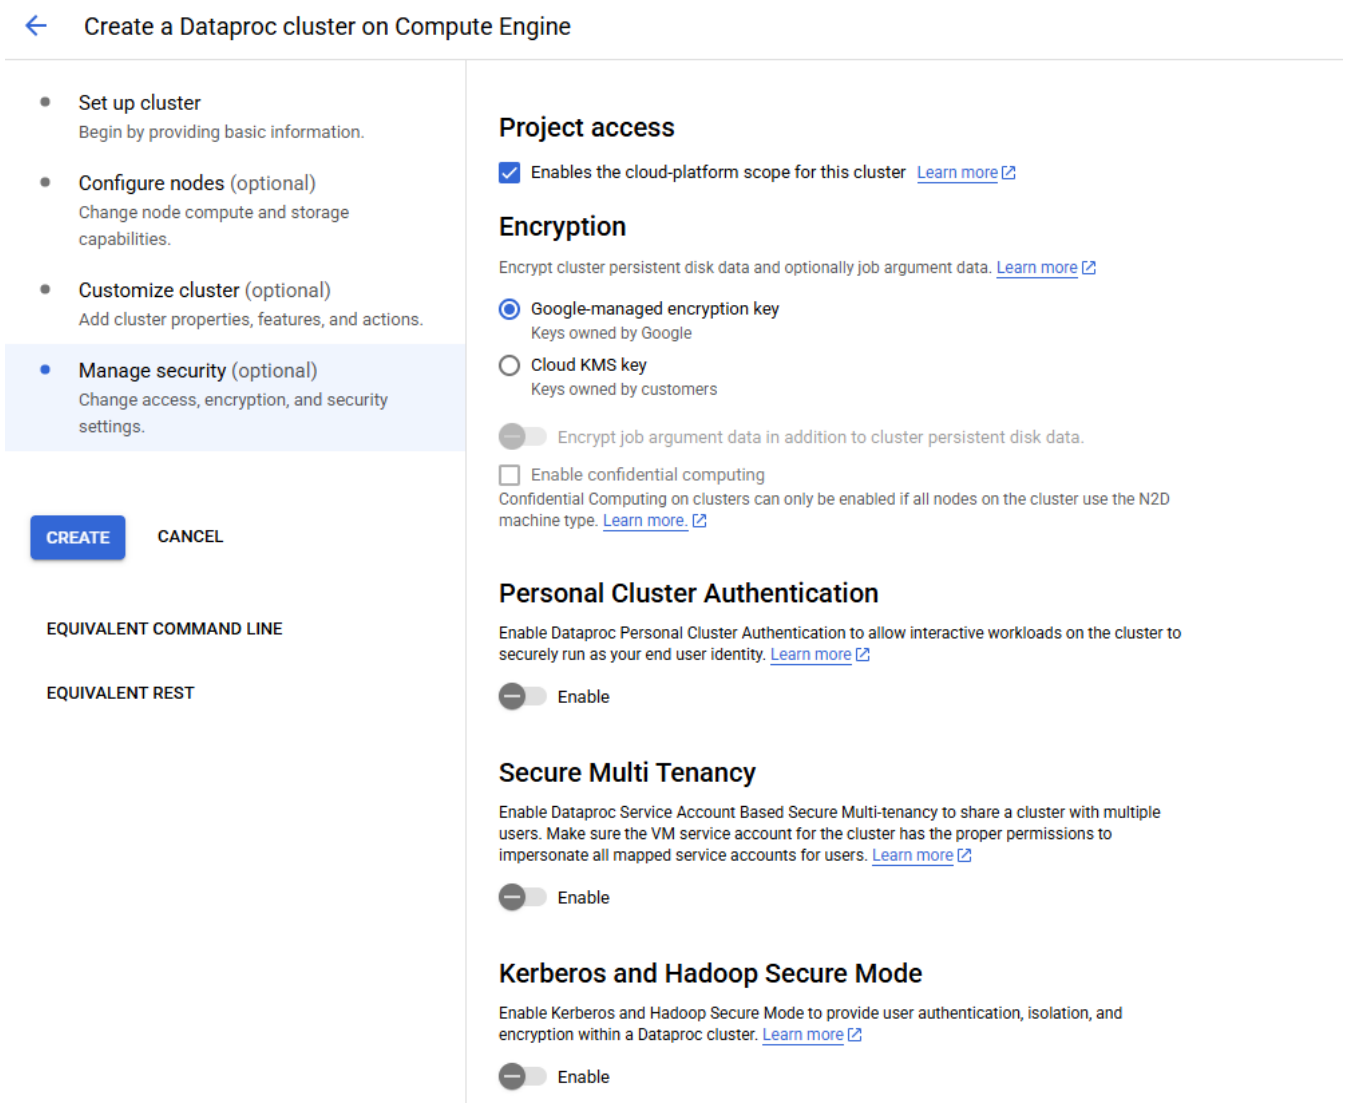

Create a Dataproc Cluster